Exploring Datastore Access Latency across AWS Compute Services

Bachelor's Thesis

Author

Bhupendra Singh

464877

b.singh@campus.tu-berlin.de

Advisor

Trever Schirmer

Examiners

Prof. Dr.-Ing. David Bermbach

Prof. Dr. habil. Odej Kao

Exploring Datastore Access Latency across AWS Compute Services

Bachelor's Thesis

Submitted by:
Bhupendra Singh
464877
b.singh@campus.tu-berlin.de

Technische Universität Berlin
Fakultät Elektrotechnik und Informatik
Fachgebiet Scalable Software Systems

2024

Hiermit versichere ich, dass ich die vorliegende Arbeit eigenständig ohne Hilfe Dritter und ausschließlich unter Verwendung der aufgeführten Quellen und Hilfsmittel angefertigt habe. Alle Stellen die den benutzten Quellen und Hilfsmitteln unverändert oder sinngemäß entnommen sind, habe ich als solche kenntlich gemacht.

Sofern generische KI-Tools verwendet wurden, habe ich Produktnamen, Hersteller, die jeweils verwendete Softwareversion und die jeweiligen Einsatzzwecke (z. B. sprachliche Überprüfung und Verbesserung der Texte, systematische Recherche) benannt. Ich verantworte die Auswahl, die Übernahme und sämtliche Ergebnisse des von mir verwendeten KI-generierten Outputs vollumfänglich selbst.

Die Satzung zur Sicherung guter wissenschaftlicher Praxis an der TU Berlin vom 8. März 2017* habe ich zur Kenntnis genommen.

Ich erkläre weiterhin, dass ich die Arbeit in gleicher oder ähnlicher Form noch keiner anderen Prüfungsbehörde vorgelegt habe.

(Unterschrift) Bhupendra Singh, Berlin, 3. Dezember 2024

---

*https://www.static.tu.berlin/fileadmin/www/10000060/FSC/Promotion___Habilitation/ Dokumente/Grundsaetze_gute_wissenschaftliche_Praxis_2017.pdf

Abstract

Cloud computing has grown rapidly, driving the demand for scalable and flexible infrastructure solutions. Amazon Web Services (AWS) leads the market, offering diverse compute and datastore services for various applications. Despite their importance, empirical data on latency dynamics between AWS compute and datastore services remains scarce. This thesis addresses this knowledge gap by systematically benchmarking access latency between EC2 and Lambda paired with RDS, DynamoDB, and S3. We evaluate these service pairs under constant and burst workloads. The results show that EC2 pairs consistently achieve lower mean latency than Lambda pairs. This study acknowledges certain limitations, including constraints imposed by AWS Free Tier, the focus on a single region, and the exclusion of mixed workload scenarios. This work provides foundational insights for designing latency-sensitive cloud applications. It emphasizes the need for further research on multi-region setups, advanced configurations, and mixed workload scenarios to deepen the understanding of AWS service interactions.

Kurzfassung

Cloud Computing hat sich schnell entwickelt und führt zu einer steigenden Nachfrage nach skalierbaren und flexiblen Infrastruktur-Lösungen. Amazon Web Services (AWS) führt den Markt an und bietet eine Vielzahl von Compute- und Datenspeicher-Diensten für unterschiedlichste Anwendungen. Trotz ihrer Bedeutung fehlen jedoch empirische Daten zu den Latenz-Dynamiken zwischen den Compute- und Datenspeicher-Diensten von AWS. Diese Arbeit schließt diese Wissenslücke, indem sie die Zugriffs-Latenz zwischen EC2 und Lambda in Kombination mit RDS, DynamoDB und S3 systematisch benchmarkt. Diese Dienst-Paare werden unter konstanten und burstigen Arbeitslasten evaluiert. Die Ergebnisse zeigen, dass EC2-Paare durchgehend eine niedrigere durchschnittliche Latenz aufweisen als Lambda-Paare. Die Arbeit erkennt bestimmte Einschränkungen an, darunter die Begrenzungen durch das AWS Free Tier, die Fokussierung auf eine einzelne Region und die Ausschließung gemischter Arbeitslastszenarien. Diese Arbeit liefert grundlegende Erkenntnisse für das Design latenzempfindlicher Cloud-Anwendungen und betont die Notwendigkeit weiterer Forschung zu Multi-Region-Setups, erweiterten Konfigurationen und gemischten Arbeitslastszenarien, um das Verständnis der Interaktionen zwischen AWS-Diensten zu vertiefen.

Contents

1 Introduction

Cloud computing delivers on-demand resources like servers, storage, and databases via the Internet, eliminating the need for physical infrastructure and enabling flexibility and scalability. Its adoption has surged in the past decade and Amazon Web Services (AWS), has emerged as a leader and pioneer of various cloud models. AWS maintains over 30% market share driven by its innovation and broad service portfolio [<**empty citation**>].

Applications in cloud environments typically consist of compute and datastore service instances that interact frequently, therefore latency between these services significantly influence overall performance and user experience of the application [<**empty citation**>]. AWS provides performance insights for its datastore services from both the service-side and client-side perspectives. The service-side perspective evaluates how efficiently the datastore processes queries internally, while the client-side perspective focuses on the time taken for a client to receive a response after submitting a query. However, it remains unclear how client-side latency varies across different AWS compute services.

Prominent AWS compute services include (1) Elastic Cloud Compute (EC2), offers customizable virtual machines for consistent performance and long-running workloads, and (2) Lambda, a serverless, event-driven service ideal for short tasks and automatic scaling. Key datastore services are (1) Relational Database Service (RDS), provides managed SQL databases for structured data, (2) DynamoDB, a serverless NoSQL database for high-scale, low-latency use cases, and (3) Simple Storage Service (S3), a distributed object storage service for unstructured data [<**empty citation**>].

In this thesis, we benchmark access latency between the above AWS compute and datastore services to examine how the choice of compute service affects latency dynamics with datastores, such as whether an EC2-RDS pair outperforms a Lambda-RDS pair. The results consistently show that EC2-based pairs outperform Lambda-based pairs by small margins. These findings can guide cloud and application developers in making informed deployment decisions, especially for latency-sensitive applications.

We therefore make the following contributions in this thesis:

1. We address the identified knowledge gap by proposing a benchmark in Section 3.

2. We evaluate the proposed approach through experimentation in Section 4.

3. We present and interpret our results in Section 5.

4. We discuss our observations, possible causes, and limitations in Section 6.

2 Background

Access latency generally means time it takes for data to reach from one point to another, also commonly referred as latency or Round-Trip Time (RTT). In the context of this thesis, it specifically refers to the time a compute service instance requires to obtain a response from datastore. It is a well known fact, that even small amounts of delay, even on microseconds scale, may lead to negative impact on the performance of an application [<empty citation>], which implies the importance of latency. From a cloud user's perspective, latency between cloud services can be influenced by certain controllable factors, such as service configuration, choice of region and availability zone, the proximity of resources, selecting instance types or service tiers optimized for low latency, network configuration, caching strategies, and application code optimization. However, underlying cloud infrastructure, especially network conditions remain uncontrollable.

Benchmarking carries different definitions across domains, however in this thesis, it specifically refers to systematically evaluating and comparing access latency between AWS compute and datastore services. Cloud benchmarking is a well-researched field, therefore the following subsections summarize the key elements and requirements of cloud benchmarking, drawing from existing literature [BWT17; Bin+09; Coo+10; Fol+13]. Additionally, we define relevant terms to provide clarity for subsequent sections.

2.1 Elements of Cloud Benchmarking

Cloud benchmarking typically involves a System Under Test (SUT) and a workload generator. The SUT may represent a single cloud service or an entire application containing components of interest. The workload generator applies artificial load on the SUT while tracking the metric of interest. In this thesis, a AWS service pair consisting of a compute instance and a datastore constitutes the SUT. However, when the compute instance itself generates load on the datastore, the tool or executable responsible for load generation is considered the workload generator.

Workload generation is a critical element of cloud benchmarking, as it directly impacts SUT. Two common workload models are open and closed models. In a closed model, a fixed number of concurrent threads independently execute a predefined sequence of tasks iteratively. In contrast, the open model specifies a rate of arrival, such as request per second. The closed model has a fundamental limitation: it ties load generation to task completion. New tasks are only scheduled once a thread completes the previous one, allowing the SUT to self-regulate. If the SUT slows or stalls, incoming load decreases, enabling recovery and potentially skewed results.

A benchmarking run refers to a single timed execution of a predefined set of tests or workloads to evaluate the performance of a SUT under specific conditions. It involves provisioning the resources, initiating the test, collecting relevant performance data, and storing the results for further analysis.

2.2 Requirements of Cloud Benchmarking

Cloud benchmarking requires several general considerations: **Relevance** ensures that the benchmarks reflect real-world scenarios and test conditions that align with the application needs. **Repeatability** ensures that results can be consistently reproduced under similar conditions, enhancing the reliability of the findings. **Fairness** is vital to ensure that all SUTs are compared equitably, with identical configurations, resources, and network

conditions to avoid bias. **Affordability** focuses on minimizing the costs of conducting the benchmark, particularly when using cloud resources with strict budget constraints. Addtionally, **simplicity** is important for creating benchmarks that are easy to understand, interpret, and build trust.

To meet advanced cloud benchmarking requirements, benchmarks should account for **failure scenarios**, testing how services handle failure conditions. They must also consider the **geo-distribution** of both measurement clients and the SUT, reflecting real-world, geographically dispersed applications. Additionally, benchmarks should provoke **stress situations** to measure qualities like **scalability and elasticity**, using variable load patterns to test service limits. The design must ensure **detailed data capture**, avoiding reliance solely on aggregate values, and should track resource consumption and costs across various components using monitoring tools. Finally, benchmarks should be **long-running** and executed across different times of day and days of the week to capture stabilized behavior, short-term effects, and seasonal variations [BWT17].

Trade-offs between requirements are inevitable, as some are inherently conflicting. For example, prioritizing relevance may lead to complex applications that reduce simplicity, while long-running experiments can increase costs, affecting affordability. Therefore, it is essential to carefully consider all requirements and make conscious trade-offs based on the benchmarking goals and scope [BWT17].

2.3 Terms and Definitions

k6[1] k6 is an open-source and extensible workload generation tool. It uses JavaScript-based test scripts to define workloads, specifying parameters such as duration, arrival rate (or concurrent threads), load patterns, and operations on the SUT. It is written in the golang with embedded JavaScript engine[2].

VPC –> TODO <–

Lambda Lambda is a serverless compute service that automatically runs code in response to events, eliminating the need to manage underlying infrastructure. It scales automatically based on incoming requests, with users only charged for execution time. It introduces cold starts, where a new execution environment is initialized when the function is invoked after being idle, causing slight latency. To reduce this, provisioned concurrency keeps a predefined number of instances ready to handle requests, while reserved concurrency limits the maximum concurrent executions per function.

DynamoDB DynamoDB is a fully managed NoSQL database service designed for low-latency, high-throughput applications. It supports the provisioned capacity mode, where users allocate Read Capacity Units (RCUs) and Write Capacity Units (WCUs) to define the number of reads and writes per second the database can handle. An RCU represents one strongly consistent read per second for an item up to 4 KB, while a WCU allows one write per second for an item up to 1 KB.

---

[1]**https://k6.io/open-source/**.
[2]**https://github.com/grafana/k6**.

3   Benchmark Design

In this section, we present our approach to benchmark AWS compute and datastore service pairs, while addressing requirements and challenges discussed in Section 2. Affordability is a major requirement in this thesis, as the whole experiment is conducted within the limits of AWS Free Tier[3]. This influences design and implementation decisions while necessitating trade-offs with other requirements.

Sections 1 and 2 highlight the importance of access latency in applications and it is the primary metric of interest in this thesis. To isolate access latency, we focus on read-only operations between compute and datastore services, to avoid locking and transactional overheads in writes. Queries are kept simple, database tables, and object files minimal to reduce querying and data transmission overhead. While mixed read-write workloads, larger databases, and complex queries reflect real-world scenarios, simple read-only benchmarks can provide initial baseline latency performance. Additionally, focusing on reads is cost-effective, as services like S3 offer significantly more free-tier reads compared to writes.

We evaluate and compare six pairs of prominent AWS compute and datastore services for access latency: (1) EC2-RDS, (2) EC2-DynamoDB, (3) EC2-S3, (4) Lambda-RDS, (5) Lambda-DynamoDB, and (6) Lambda-S3. Each benchmarking run involves two primary components: (1) a pre-loaded datastore serving as the System Under Test (SUT) and (2) a compute instance which generates workload on the SUT and collects access latency data. For simplicity, we refer to the datastore as SUT and the compute instance as workload generator in this thesis. However, in our context, a compute instance and a datastore together constitute one SUT, as we are not particularly testing the datastores.

To address relevance, we employ open workload model, due to the reasons discussed in Section 2.1 and evaluate each pair under two distinct workload types: constant and burst. Under constant workload, the reads per second (RPS) remain stable throughout the benchmarking run. A burst workload, on the other hand, simulates sudden spikes in RPS, reflecting real-world scenarios where workload can be unpredictable. This allows us to observe how each pair performs under steady load and how they respond to fluctuating, high-demand conditions. To address repeatability we evaluate each pair under both workload types twice, resulting in total of 24 runs.

As discussed in Section 2.3, the fluctuating underlying conditions of the cloud can impact performance. To address this challenge, we extend the durations of benchmarking runs as much as possible and repeat them at different times and days, while balancing affordability and relevance. Fairness is ensured by maintaining identical workloads, resource allocation, network conditions, and datastore configurations across service pairs.

---

[3]**https://aws.amazon.com/free/**.

4 Benchmark Implementation

This section outlines the implementation details of the benchmark, focusing on three main aspects: (1) service instance configurations, (2) workload generation, and (3) data analysis.

For EC2-Pairs, the benchmarking process is straightforward, with the EC2 instance generating workload on datastores using k6 and storing the collected data locally. For Lambda-Pairs, each Lambda invocation corresponds to a single read operation on the datastore. To invoke the Lambda function programmatically and collect results, an additional EC2 instance, termed as "Lambda-Helper," is provisioned with k6 and the necessary test scripts, as shown in **??**.

All benchmarking runs are executed in the *eu-central-1* AWS region. All components of a run are provisioned and de-provisioned using Infrastructure-as-Code combined with shell scripting. This approach ensures a fresh environment for every run and minimizes the potential for human error. Upon successful completion of a run, the collected data is exported to a dedicated S3 bucket. Additionally, to keep network conditions and resource allocation identical, no benchmarking runs are executed simultaneously.

−> Diagram here <−

4.1 Resource Configuration

RDS: A MySQL 8.0 instance on a t3.micro virtual machine (VM) with 1GiB RAM, 2 vCPUs, and 5GB storage. The instance is located in the eu-central-1a availability zone, with multi-AZ failover disabled to maintain a single-zone configuration. It hosts a single database containing a table with 1000 rows and four columns.

DynamoDB: A DynamoDB table with 10 items, read capacity of 25 Read Capacity Units (RCUs), write capacity of 5 Write Capacity Units (WCUs), and configured with *Provisioned* capacity mode.

S3: A single S3 bucket with a 20-byte text file.

EC2: An EC2 instance of t2.micro type with 1GiB RAM, and 1vCPU, located in the eu-central-1a availability zone. It is operated by Ubuntu Server 24.04 (HVM-based) and is equipped with the k6 and corresponding test scripts.

Lambda: A non-VPC Node.js 20 function with 1.65GiB RAM, 1vCPU, and reserved concurrency set to 990. The configuration of Lambda-Helper is identical to that of EC2. CPU-Boost for cold starts?

4.2 Workload Generation

With following workload configuration we specifically address relevance, affordability, repeatability, and cloud performance fluctuations. However, for relevance, there is no standard rate or load pattern at which, for instance, a backend server reads from a datastore, and depends highly on the application's use-case and demand. Instead, we focus on extending the run duration to a point where it can be repeated twice while remaining within AWS Free Tier. Time limitations are also considered.

Constant Workload:

- **RDS and DynamoDB**: 2 RPS for 14 hours.

- **S3**: 0.2 RPS (1 read every 5 seconds) for 3 hours.

Burst Workload:

- **RDS and DynamoDB**: The run starts with a baseline of 2 RPS for the first 2 hours, followed by 30 load spikes. Each spike lasts 3 minutes, with RPS increasing linearly from 2 to 20 in 90 seconds and then decreasing back to the baseline of 2 in 90 seconds. Between spikes, the load remains at 2 RPS for 3 minutes. The run concludes with an additional 2 hours at the 2 RPS, resulting in a total test duration of approximately 7 hours.

- **S3**: The run starts with a baseline of 0.2 RPS for the first 20 minutes, followed by 6 load spikes. Each spike lasts 30 seconds, with RPS increasing linearly from 0.2 to 16 in 15 seconds and then decreasing back to the baseline of 0.2 in 15 seconds. Between spikes, the load remains at 0.2 RPS for 6 minutes. The run concludes with an additional 15 minutes at the 0.2 RPS, resulting in a total test duration of approximately 70 minutes.

As noted, the monthly limits for S3 are relatively low, specifically 20000 GET requests per month, which resulted in shorter runs with S3.

4.3   Data Analysis

As discussed, for each targeted datastore service, we compare latency performance across EC2 and Lambda. Mean values alone cannot fully capture the behavior of access latency throughout the experiment, so we use bar plots for aggregated metrics and time-series graphs to provide a comprehensive view of the data.

In our context, outliers in latency values may indicate potential network issues between services and can distort analysis, particularly when calculating averages [BWT17]. To address this, we identify and remove outliers using the 3-sigma rule, which defines data points beyond three standard deviations from the mean as outliers [HKP11]. However, the number and extent of these outliers provide valuable insights into a service pair's performance and reliability. Therefore, we also analyze and report outliers to understand their implications and highlight variations between pairs.

5   Benchmark Results

In this section, we present and compare the findings from the data analysis for each targeted datastore, reporting values on a millisecond (ms) scale. Throughout the benchmarking runs, no CPU performance issues are observed on the client side (workload generator), ensuring that latency measurements remain unaffected by client limitations. Furthermore, the error rate consistently registers zero across all runs, confirming that every read request is successfully executed.

EC2-RDS vs. Lambda-RDS   Figure 5.1.1 and 5.1.2 show that EC2-RDS outperforms Lambda-RDS by approximately 51% under constant and 56% under burst workload. EC2-RDS shows 41% and 43% lower standard deviation under constant and burst workload, respectively. Maximum latency values of EC2-RDS are 82% and 56% lower under constant and burst workload, respectively. Figure 5.2.1 and 5.2.2 show that Lambda-RDS latency often experiences abrupt increases or decreases and persists at such levels for several hours, but in contrast EC2-RDS latency remains comparatively stable. The second runs of EC2-RDS under both workload patterns outperform the initial runs. Across both runs and workloads, EC2-RDS records 962 outliers with latencies between $5ms$ and $464ms$, while Lambda-RDS records 1879 outliers, ranging from $11ms$ to $818ms$.

EC2-DynamoDB vs. Lambda-DynamoDB   Figure 5.1.3 and 5.1.4 show that EC2-DynamoDB outperforms Lambda-DynamoDB by approximately 14% under constant and 16% under burst workload. Lambda-DynamoDB shows 21% and 7% lower standard deviation under constant and burst workload, respectively. Maximum latency values are identical across both pairs and modes. Although Lambda-DynamoDB shows lower latency variance, Figure 5.2.3 and 5.2.4 reveal that Lambda-DynamoDB experiences sudden temporary latency spikes under both workloads. These spikes reach peaks significantly higher than those observed with EC2-DynamoDB. Similar to EC2-RDS, the second runs of EC2-DynamoDB under both workloads outperform the initial runs, especially under burst workload. Across both runs and workloads, EC2-DynamoDB records 3956 outliers with latencies between $7ms$ and $100ms$, while Lambda-DynamoDB records 3073 outliers, ranging from $8ms$ to $209ms$.

EC2-S3 vs. Lambda-S3   Figure 5.1.5 and 5.1.6 show that EC2-S3 outperforms Lambda-S3 by approximately 9% under both constant and burst workload. Lambda-S3 shows 20% lower standard deviation under constant but 15% higher under burst workload. Maximum latency value of Lambda-S3 is 32% lower under constant but 11% higher under burst workload. Unlike RDS and DynamoDB, Figure 5.2.5 and 5.2.6 show that Lambda-S3 latency does not experience temporal performance shifts or large spikes. Additionally, for EC2-S3, the second runs under both workload patterns do not show better performance compared to the initial runs, a trend observed with RDS and DynamoDB. Across both runs and workloads, EC2-S3 records 128 outliers with latencies between $22ms$ and $834ms$, while Lambda-S3 records 148 outliers, ranging from $26ms$ to $390ms$.

Summary   The benchmark results indicate that RDS, DynamoDB, and S3 paired with EC2 generally outperform those paired with Lambda in terms of mean latency. However, the absolute difference between EC2- and Lambda-based pairs mean latencies remains minimal, consistently under $1.34ms$ across all experiments. Standard deviation and outliers, on the other hand, do not show a consistent pattern across the pairs.

(5.1.1) Constant Workload on RDS

(5.1.2) Burst Workload on RDS

(5.1.3) Constant Workload on DynamoDB

(5.1.4) Burst Workload on DynamoDB

(5.1.5) Constant Workload on S3

(5.1.6) Burst Workload on S3

Figure 5.1: Aggregation of Latency Measurements

(5.2.1) Constant Workload on RDS

(5.2.2) Burst Workload on RDS

(5.2.3) Constant Workload on DynamoDB

(5.2.4) Burst Workload on DynamoDB

(5.2.5) Constant Workload on S3

(5.2.6) Burst Workload on S3

Figure 5.2: Time-Series Representation of Latency Measurements

6   Discussion

Section 5 shows that datastore access latency is influenced by the choice of compute service, with EC2 consistently outperforming Lambda across all experiments. This suggests an inherent variance and overhead associated with Lambda. In this section, we explore possible reasons for these observations and also assess our approach and its limitations.

In Figure 5.2, we observe temporal shifts in the latency performance, even under constant workload. Temporal performance variability in serverless platforms, including AWS Lambda, is well-documented [Eis+22; GF21; Sch+23]. Ginzburg et al. [GF21] report significant performance shifts in AWS Lambda during a one-week continuous benchmark, even suggesting that these variations could be exploited to reduce costs. It is possible that our benchmarking runs, specifically Lambda-RDS (Figure 5.2.1 and Figure 5.2.2) and Lambda-Dynamo (Figure 5.2.3 and Figure 5.2.4), were affected by such temporal shifts, impacting the results. Similar is also applicable for EC2 [IYE11; SDQR10], however, Dancheva et al. [DAB24] suggest that recent advances are able to improve networking performance in EC2.

Lambda functions run inside microVMs based on Firecracker[4] hypervisor, deployed on bare-metal, multi-tenant EC2 instances referred as Lambda-Workers. Unlike standard EC2 instances based on the AWS Nitro[5] hypervisor, Lambda-Workers add Firecracker as an additional abstraction layer. Communication between a microVM and Firecracker occurs via optimized virtio[6] interface [Bro+23], leading to overall networking overhead of approximately $0.06ms$ under controlled, isolated conditions [Rep]. In production, however, Lambda-Workers may host hundreds or thousands of microVMs [Aga+20], potentially increasing this overhead due to induced network noise and resource contention [DS+22; Wan+18]. Agache et al. [Aga+20] also highlight Firecracker's comparatively lower networking performance against other hypervisors. Furthermore, Lambda-Workers are hosted within a network-isolated VPC managed by Lambda in the service accounts inaccessible to customers [TBT22], leaving their precise network configuration unclear.

Limitations   In Section 2, we outlined general aspects and requirements of cloud benchmarking, presented our approach in Section 3, and validated it through experimentation. However, affordability emerged as a significant limitation. AWS Free Tier constraints restricted experiment duration and scale, limiting long-term performance trend observations. The use of low-tier configurations, such as t3.micro instances and no provisioned concurrency for Lambda, hindered evaluation under realistic production conditions. Temporal variability in cloud performance posed another challenge, with short-duration benchmarks, particularly for S3, possibly failing to capture the real empirical distribution of access latency. Additionally, the absence of multi-region setups excluded analysis of geo-distributed latency behaviors. Lastly, latency measurements were limited to $1ms$ precision, whereas observed differences often occurred at the microsecond scale, suggesting that higher-precision timestamps with $1\mu s$ accuracy could yield more accurate insights.

---

[4]**https://firecracker-microvm.github.io/**.
[5]**https://aws.amazon.com/ec2/nitro/**.
[6]**https://libvirt.org/**.

7 Related Work

$-> $ TODO $<-$

8 Conclusion

Cloud computing has experienced rapid growth, becoming integral to industries seeking scalable and flexible solutions. AWS has emerged as a leader, offering diverse services for computing and storage. Despite its significance, limited empirical data exists on the latency characteristics of AWS compute-datastore interactions, necessitating this study to bridge the knowledge gap.

This thesis investigated the latency performance between AWS compute and datastore service pairs, specifically EC2 and Lambda paired with RDS, DynamoDB, and S3. A systematically conducted benchmark evaluated each pair under constant and burst workloads. Results indicated that EC2 pairs consistently demonstrated lower latency compared to Lambda pairs.

The study highlighted the influence of computing service choices on latency performance and provided empirical data to support architectural decisions in cloud environments. Despite limitations, such as the constraints of AWS Free Tier and the exclusion of mixed workload scenarios, the findings offer foundational insights into the dynamics of service interactions in AWS.

Future work can expand on these results by exploring higher-tier configurations, multi-region setups, and mixed read-write workloads to provide a more comprehensive understanding of cloud service latency characteristics.

References

[Aga+20] Alexandru Agache, Marc Brooker, Alexandra Iordache, Anthony Liguori, Rolf Neugebauer, Phil Piwonka, and Diana-Maria Popa. "Firecracker: Lightweight Virtualization for Serverless Applications". In: *17th USENIX Symposium on Networked Systems Design and Implementation (NSDI 20)*. Santa Clara, CA: USENIX Association, Feb. 2020, pp. 419–434. ISBN: 978-1-939133-13-7.

[BWT17] David Bermbach, Erik Wittern, and Stefan Tai. *Cloud Service Benchmarking: Measuring Quality of Cloud Services from a Client Perspective*. Cham, Switzerland: Springer, 2017. ISBN: 978-3-319-85672-8.

[Bin+09] C. Binnig, D. Kossmann, T. Kraska, and S. Loesing. "How is the weather tomorrow? Towards a benchmark for the cloud". In: *Proceedings of the Second International Workshop on Testing Database Systems*. 2009, pp. 1–6.

[Bro+23] Marc Brooker, Mike Danilov, Chris Greenwood, and Phil Piwonka. "On-demand Container Loading in AWS Lambda". In: *2023 USENIX Annual Technical Conference (USENIX ATC 23)*. Boston, MA: USENIX Association, July 2023, pp. 315–328. ISBN: 978-1-939133-35-9.

[Coo+10] Brian F. Cooper, Adam Silberstein, Erwin Tam, Raghu Ramakrishnan, and Russell Sears. "Benchmarking Cloud Serving Systems with YCSB". In: *Proceedings of the 1st ACM Symposium on Cloud Computing* (Indianapolis, Indiana, USA). SoCC '10. New York, NY, USA: Association for Computing Machinery, 2010, pp. 143–154. DOI: `10.1145/1807128.1807152`.

[DAB24] Tamara Dancheva, Unai Alonso, and Michael Barton. "Cloud benchmarking and performance analysis of an HPC application in Amazon EC2". In: *Cluster Computing* 27.2 (2024), pp. 2273–2290.

[DS+22] Daniele De Sensi, Tiziano De Matteis, Konstantin Taranov, Salvatore Di Girolamo, Tobias Rahn, and Torsten Hoefler. "Noise in the Clouds: Influence of Network Performance Variability on Application Scalability". In: *Proc. ACM Meas. Anal. Comput. Syst.* 6.3 (Dec. 2022). DOI: `10.1145/3570609`.

[Eis+22] Simon Eismann, Diego Elias Costa, Lizhi Liao, Cor-Paul Bezemer, Weiyi Shang, André van Hoorn, and Samuel Kounev. "A case study on the stability of performance tests for serverless applications". In: *Journal of Systems and Software* 189 (2022), p. 111294. ISSN: 0164-1212. DOI: `https://doi.org/10.1016/j.jss.2022.111294`.

[Fol+13] Enno Folkerts, Alexander Alexandrov, Kai Sachs, Alexandru Iosup, Volker Markl, and Cafer Tosun. "Benchmarking in the cloud: What it should, can, and cannot be". In: *Selected Topics in Performance Evaluation and Benchmarking: 4th TPC Technology Conference, TPCTC 2012, Istanbul, Turkey, August 27, 2012, Revised Selected Papers 4*. Springer. 2013, pp. 173–188.

[GF21] Samuel Ginzburg and Michael J. Freedman. "Serverless Isn't Server-Less: Measuring and Exploiting Resource Variability on Cloud FaaS Platforms". In: WoSC '20. Delft, Netherlands: Association for Computing Machinery, 2021, 43–48. ISBN: 9781450382045. DOI: `10.1145/3429880.3430099`.

[HKP11] J. Han, M. Kamber, and J. Pei. *Data Mining: Concepts and Techniques*. The Morgan Kaufmann Series in Data Management Systems. Morgan Kaufmann, 2011. ISBN: 9780123814807.

[IYE11] Alexandru Iosup, Nezih Yigitbasi, and Dick Epema. "On the performance variability of production cloud services". In: *2011 11th IEEE/ACM International Symposium on Cluster, Cloud and Grid Computing*. IEEE. 2011, pp. 104–113.

[Rep] *Firecracker: Secure and fast microVMs for serverless computing.* Amazon Web Services. Dec. 2024. URL: https://github.com/firecracker-microvm/firecracker/.

[SDQR10] Jörg Schad, Jens Dittrich, and Jorge-Arnulfo Quiané-Ruiz. "Runtime measurements in the cloud: observing, analyzing, and reducing variance". In: *Proc. VLDB Endow.* 3.1–2 (Sept. 2010), 460–471. ISSN: 2150-8097. DOI: 10.14778/1920841.1920902.

[Sch+23] Trever Schirmer, Nils Japke, Sofia Greten, Tobias Pfandzelter, and David Bermbach. "The Night Shift: Understanding Performance Variability of Cloud Serverless Platforms". In: *Proceedings of the 1st Workshop on Serverless Systems, Applications and Methodologies.* SESAME '23. Rome, Italy: Association for Computing Machinery, 2023, 27–33. ISBN: 9798400701856. DOI: 10.1145/3592533.3592808.

[TBT22] Mayank Thakkar, Marc Brooker, and AWS Lambda Security Team. *Security Overview of AWS Lambda.* Amazon Web Services. Dec. 2022. URL: https://docs.aws.amazon.com/whitepapers/latest/security-overview-aws-lambda/.

[Wan+18] Liang Wang, Mengyuan Li, Yinqian Zhang, Thomas Ristenpart, and Michael Swift. "Peeking Behind the Curtains of Serverless Platforms". In: *2018 USENIX Annual Technical Conference (USENIX ATC 18).* Boston, MA: USENIX Association, July 2018, pp. 133–146. ISBN: ISBN 978-1-939133-01-4.